

# Implementation of DIAL-MPC for a Robotic Arm Lift Task in ManiSkill3

April 2, 2025

## Abstract

This document presents an outline for adapting and implementing DIAL-MPC (Diffusion-Inspired Annealing for Model Predictive Control) to solve a robotic arm lift task in the ManiSkill3 environment. The approach leverages GPU-based parallel rollouts and a diffusion-style annealing framework to perform online trajectory optimization without requiring offline training.

## Contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
<b>2</b>	<b>Environment Setup and Parallel Simulation</b>	<b>2</b>
2.1	ManiSkill3 Environment Configuration . . . . .	2
2.2	Creating Parallel Environments . . . . .	2
<b>3</b>	<b>Cost Function Design and Reward Integration</b>	<b>2</b>
3.1	Analyzing Built-in Rewards . . . . .	2
3.2	Converting Rewards to Costs . . . . .	3
3.3	Accumulating Cost . . . . .	3
<b>4</b>	<b>Core Implementation of DIAL-MPC</b>	<b>3</b>
4.1	Initialization . . . . .	3
4.2	Outer Loop: Diffusion-Inspired Annealing . . . . .	3
4.3	Inner Loop: Action-Level Annealing . . . . .	3
4.4	Parallel Rollouts and Weighted Update . . . . .	3
<b>5</b>	<b>Control Loop and Execution</b>	<b>4</b>
5.1	Action Execution . . . . .	4
5.2	Sliding Window Mechanism . . . . .	4
5.3	State Synchronization . . . . .	4
<b>6</b>	<b>Parameter Settings and Practical Tips</b>	<b>4</b>
6.1	Control Frequency and Horizon . . . . .	4
6.2	Parallel Sampling Quantity . . . . .	4
6.3	Annealing Schedule and Temperature Parameter . . . . .	4
6.4	Cost Function Tuning . . . . .	4

<b>7</b>	<b>Testing and Debugging</b>	<b>4</b>
7.1	Offline Simulation Testing . . . . .	4
7.2	Parameter Tuning . . . . .	5
7.3	Benchmark Comparison . . . . .	5
7.4	Real-World Deployment . . . . .	5

## 1 Introduction

DIAL-MPC [1] is a sampling-based Model Predictive Control framework that combines large-scale parallel rollouts with a diffusion-inspired dual-loop annealing strategy. Originally developed for legged robot control, we aim to adapt it here for a robotic manipulator lift task in ManiSkill3.

- No offline training needed, broad global search plus local convergence, real-time optimization at each control timestep.
- Handling high-dimensional manipulator state, contact forces, environment constraints, and ensuring stable real-time performance.

## 2 Environment Setup and Parallel Simulation

### 2.1 ManiSkill3 Environment Configuration

- Select a Lift scenario in ManiSkill3 (environment ID: PickCube-v1) that involves grasping and lifting an object with a Panda robot arm.
- Ensure that relevant packages (SAPIEN, GPU drivers) are correctly installed, and GPU-based parallel simulation is enabled.
- Verify that each environment provides consistent observations (joint states, object position, etc.).

### 2.2 Creating Parallel Environments

- Instantiate  $N$  parallel instances of the lift environment using ManiSkill3’s vector interfaces or custom multi-process solutions. (My GPU is 3060ti, the parallel environment should be small).
- Use identical initial states if needed, or randomize according to the task’s setup.
- Each environment must support rollouts up to the planning horizon  $H$  in parallel.

## 3 Cost Function Design and Reward Integration

### 3.1 Analyzing Built-in Rewards

- ManiSkill3 Lift tasks often provide multi-stage rewards (grasp success, stable lift, final height).
- Identify each reward term to decide how they translate into cost terms.

### 3.2 Converting Rewards to Costs

- DIAL-MPC expects a cost to minimize, but environment typically gives reward.
- Define cost as  $\text{cost} = -(\text{reward})$ ; optionally re-scale or offset if needed.

### 3.3 Accumulating Cost

- For each candidate action sequence, sum the per-step cost:

$$J(u_0, \dots, u_H) = \sum_{t=0}^H [-r_t].$$

- If needed, add extra terms (e.g., torque limit penalties, collision cost).

## 4 Core Implementation of DIAL-MPC

### 4.1 Initialization

- Define a horizon  $H$ , an initial guess  $U^{(0)} = \{u_0, \dots, u_H\}$ , and initial covariance  $\Sigma_h^{(0)}$  for sampling each time step  $h$ .
- Copy the current state from the real or main environment into the parallel environments.

### 4.2 Outer Loop: Diffusion-Inspired Annealing

- Perform  $N_{\text{outer}}$  iterations of the annealing process.
- Gradually decrease covariance from large (global exploration) to small (local refinement).
- Example: exponential schedule for  $\det(\Sigma^i)$ .

### 4.3 Inner Loop: Action-Level Annealing

- Different time steps in the horizon can have different noise scales.
- Let near-term actions be more precise (small noise), while far-future actions keep more exploration (larger noise).

### 4.4 Parallel Rollouts and Weighted Update

- Sample  $N_{\text{samples}}$  sequences around the current solution; evaluate cost by rolling out environment up to horizon  $H$  in parallel.
- Use an exponential weighting scheme:

$$w_j = \exp\left(-\frac{J_j}{\lambda}\right),$$

then update  $U^+$  via the weighted average of perturbations.

## 5 Control Loop and Execution

### 5.1 Action Execution

- Once the solution is found, apply the first action  $u_0$  to the real manipulator.
- Maintain the rest of the sequence for the next time step’s warm-start.

### 5.2 Sliding Window Mechanism

- Shift  $U \leftarrow (u_1, u_2, \dots, u_H, \text{new guess})$ .
- This ensures continuity of control actions from one MPC iteration to the next.

### 5.3 State Synchronization

- After applying  $u_0$  in the real environment, measure new state (joint angles, object pose, etc.).
- Reset parallel environments to that state for the next iteration’s rollouts.

## 6 Parameter Settings and Practical Tips

### 6.1 Control Frequency and Horizon

- A 10–20 Hz MPC loop is common for manipulator tasks.
- Horizon  $H \approx 10\text{--}20$  steps (0.5–1.0 s lookahead).

### 6.2 Parallel Sampling Quantity

- The number of samples  $N_{\text{samples}}$  per iteration depends on GPU capacity.
- Typical range: 256–2048.

### 6.3 Annealing Schedule and Temperature Parameter

- Use exponential or piecewise-linear scheduling for  $\Sigma$ .
- Adjust  $\lambda$  (temperature) to match cost magnitudes and avoid numerical instability.

### 6.4 Cost Function Tuning

- Add minor penalties for large joint velocities or excessive torque to ensure stable motion.
- If the arm exhibits jitter, increase the final annealing stage length or reduce noise.

## 7 Testing and Debugging

### 7.1 Offline Simulation Testing

- Validate that the arm can indeed lift the object in a pure simulation environment.
- Monitor the cost convergence and final success rate.

## 7.2 Parameter Tuning

- Adjust  $N_{\text{outer}}$ ,  $N_{\text{samples}}$ ,  $\Sigma$  schedules, and  $\lambda$  iteratively.
- Too large a noise leads to divergence; too small a noise leads to local minima.

## 7.3 Benchmark Comparison

- Compare with SAC algorithms in terms of success rate and policy generalization.
- Evaluate energy consumption or torque usage as additional metrics.

## 7.4 Real-World Deployment

- Account for latency, sensor noise, domain gap between real manipulator and simulation.
- Possibly learn an online correction model for friction or contact differences.

## References

- [1] H. Xue et al., *Diffusion-Inspired Annealing for Legged MPC (DIAL-MPC)*, arXiv:2409.15610, 2024.